

Asmi Yadav

+91 99931 00849 | asmiyadav2004@gmail.com | [linkedin.com/in/asmi-yadav-b0824724a](https://www.linkedin.com/in/asmi-yadav-b0824724a) | github.com/neuralasmi

Summary

Computer Science graduate (IoT & Intelligent Systems, 2026) who builds AI models and also ships them: fine-tuned transformers, RAG pipelines, and the FastAPI, Docker and Linux layer underneath. Strongest in AI/ML with backend and systems close behind, and I build fast with Claude and similar tools.

Experience

FRND — AI Content Analyst (Intern)

Jan 2026 – Feb 2026

Studied large conversational datasets to find response gaps, repetition, and inconsistent model behavior, then used those findings to clean up training data and improve chatbot quality. Built the analysis in Python with HuggingFace Transformers and scikit-learn.

Tiffinly — Community Meal Marketplace (Freelance)

Apr 2025 – Sept 2025

Built and shipped the full MVP of a marketplace connecting college students with nearby home cooks: REST API backend, meal request posting, cook matching by location, and in-app chat for price and timing. Ran a real pilot with 50 students and 10 cooks, and handled the whole thing myself end to end.

Projects

CineSearch — RAG QA Engine with Local LLM Inference

github.com/neuralasmi/CineSearch

FastAPI · LangChain · FAISS · Ollama/Llama3 · Docker Compose · SSE streaming

A RAG system for question answering over your own documents. Ingests and chunks PDFs with LangChain, embeds with sentence-transformers, stores vectors in FAISS, and streams answers over FastAPI Server-Sent Events. Llama3 runs locally through Ollama in a docker-compose setup, so nothing depends on a hosted API. Sub-second retrieval across 100+ documents.

StreamServe — FastAPI + Docker Model Serving Pipeline

github.com/neuralasmi/StreamServe

FastAPI · Docker · scikit-learn · RandomForest · REST API

A model serving pipeline built to be deployed, not demoed. RandomForest trained and exported with joblib, then served behind FastAPI with GET /health and POST /predict. Multi-stage Docker build on python:3.11-slim, training kept separate from serving.

MetaTag — BERT NER for Content Metadata

github.com/neuralasmi/MetaTag

bert-base-uncased · HuggingFace · CoNLL-2003 · token classification

Fine-tuned bert-base-uncased on CoNLL-2003 across four entity types (PER, ORG, LOC, MISC), reaching entity-level F1 ~ 0.91 on the test set. Handled word-to-subword label alignment (-100 masking) plus a modular training loop with logging, checkpointing, and sequeval evaluation.

Technical Skills

Languages:	Python, Java, C, C++, SQL, HTML/CSS
AI/ML:	PyTorch, HuggingFace Transformers, Scikit-learn, XGBoost, LangChain, FAISS, Pandas, NumPy
Backend & Systems:	Linux & command line, Docker, docker-compose, FastAPI, REST and streaming (SSE) APIs, Git
Data & MLOps:	MongoDB, PyMongo, MySQL, MLflow, Weights & Biases, Postman
AI-Assisted Dev:	Claude / Claude Code, ChatGPT — daily driver for building and debugging

Education

Manipal University Jaipur

Aug 2022 – Aug 2026

B.Tech Computer Science (IoT & Intelligent Systems) | Jaipur, Rajasthan

St. Joseph's Coed Sr. Sec. School, Bhopal

Mar 2021 – Mar 2022

CBSE 12th Boards – 91.3%

Certifications

- Mastering Data Structures and Algorithms (C/C++) — Abdul Bari
- AI Tools and ChatGPT Workshop